


A blurred background image showing a group of people, likely a diverse group, giving thumbs up. The focus is on the hands and the general positive gesture, with the faces being out of focus.

# RECOMMENDATION



# EXAMPLES

- Diapers and beer (legend?)
- Target pregnant teenager



# EXAMPLES

- Netflix Prize 2009



- Amazon Recommendation Engine

## Customers Who Bought This Item Also Bought



[30 Rock: Seasons 1-3](#)  
DVD ~ Tracy Morgan  
★★★★★ (7)  
\$60.49



[Desperate Housewives: The Complete Seasons 1-5](#)  
DVD ~ Teri Hatcher  
★★★★☆ (2)  
\$179.99

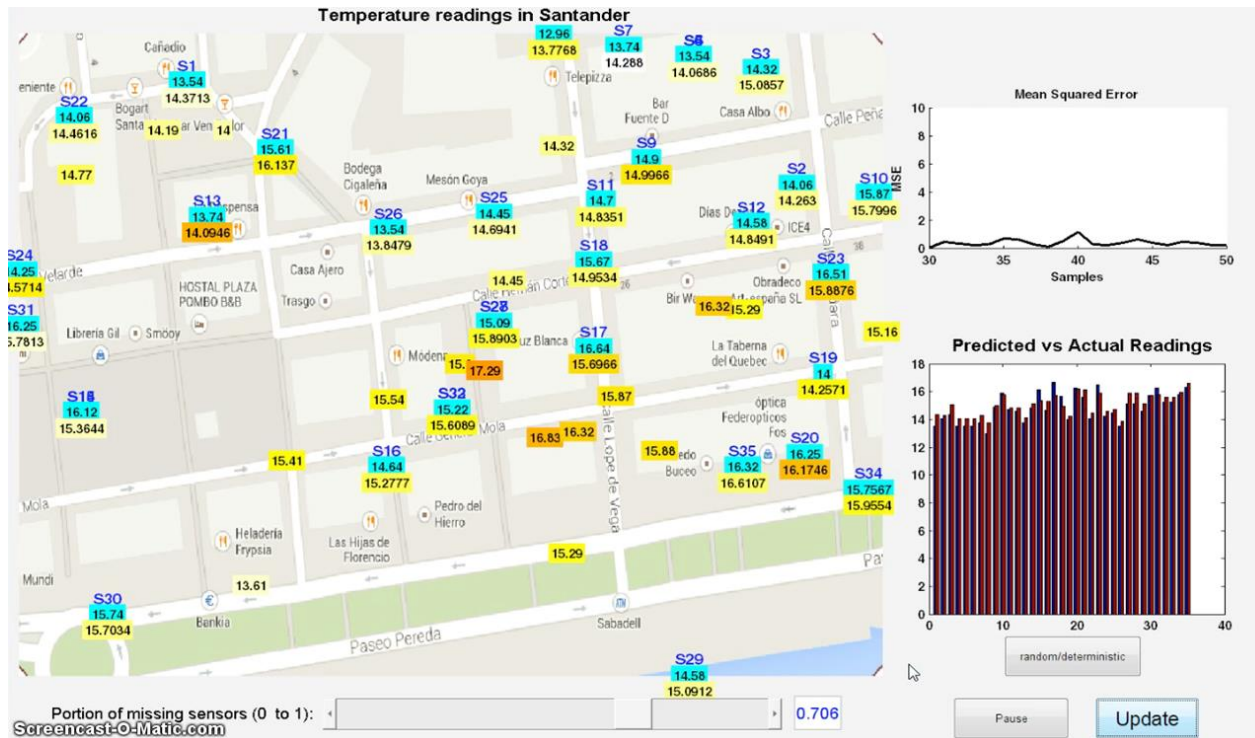


[Scrubs: The Complete Seasons 1-8](#)  
DVD ~ Zach Braff  
★★★★★ (2)  
\$148.49



# EXAMPLES

- Missing sensor data





# COLLABORATIVE FILTERING

- 400,000 users  
17,000 movies  
But only a few ratings (1%)
- User  $a$   
movie  $i$   
Rating  $Y_{ai} \in \mathbb{R}$  (e.g. 1-5)
- Goal is to predict  
unobserved ratings

$m$  movies

$n$  users

5	5						5		
		3	5	1	3	4	4		4
	4	2			2				
		5							5
4	5							4	
4							4		
5		4	5	1		4			
	4								
5				4					
5						4			
		5				5		3	

$Y_{ai}$



# COLLABORATIVE FILTERING

- Collaborative: cross-users  
Filtering: prediction
- Matrix or tensor completion problems

A tensor is a multi-dimensional array.  
e.g.  $n \times m$  2-tensor  
e.g.  $p \times q \times r$  3-tensor  
(2-tensor = matrix)

$m$  movies

5	5						5		
		3	5	1	3	4	4		4
	4	2			2				
		5							5
4	5							4	
4							4		
5		4	5	1		4			
	4								
5				4					
5						4			
		5				5		3	

$n$  users

$Y_{ai}$






# NEAREST NEIGHBOR



# K-NEAREST NEIGHBORS

## Basic Idea.

- Find **a few users**  $b_1, \dots, b_k$  (neighbors) that are similar to user  $a$
- Use information from users  $b_1, \dots, b_k$  to predict ratings of user  $a$





# DISTANCE METRIC

Measure **statistical correlation** between users  $a$  and  $b$

$$\text{sim}(a, b) = \text{corr}(a, b) = \frac{\sum_{j \in CR(a, b)} (Y_{aj} - \tilde{Y}_a)(Y_{bj} - \tilde{Y}_b)}{\sqrt{\sum_{j \in CR(a, b)} (Y_{aj} - \tilde{Y}_a)^2} \sqrt{\sum_{j \in CR(a, b)} (Y_{bj} - \tilde{Y}_b)^2}} \in [-1, 1]$$

cosine  
similarity!

where

1.  $CR(a, b)$  is the set of movies rated by both  $a$  and  $b$  (common ratings)
2.  $\tilde{Y}_a = \frac{1}{|CR(a, b)|} \sum_{j \in CR(a, b)} Y_{aj}$   
(average rating for user  $a$  among common ratings)



# PREDICTION

$$\hat{Y}_{ai} = \bar{Y}_a + \frac{\sum_{b \in kNN(a,i)} \text{sim}(a, b)(Y_{bi} - \bar{Y}_b)}{\sum_{b \in kNN(a,i)} |\text{sim}(a, b)|}$$

where

1.  $\bar{Y}_a$  is the mean rating for user  $a$
2.  $kNN(a, i)$  are the  $k$  nearest neighbors of user  $a$  that also rated movie  $i$



# DISCUSSION

$$\hat{Y}_{ai} = \bar{Y}_a + \frac{\sum_{b \in kNN(a,i)} \text{sim}(a, b)(Y_{bi} - \bar{Y}_b)}{\sum_{b \in kNN(a,i)} |\text{sim}(a, b)|}$$

- Weighted sum (could be negative weights).  
Should we consider anti-correlated neighbors?
- Good to account for bias (mean) of each user,  
but sensitive to the spread (variance) of each user.
- How do we choose the optimal  $k$ ?
  - Minimize validation error
  - Problem statement
  - Computational resources





# MATRIX FACTORIZATION



# LOW-RANK APPROXIMATIONS

## Idea.

Assume **fully observed**  $\hat{Y}$  lie in a small class  $\mathcal{H}$  of matrices.

Find matrix in  $\mathcal{H}$  that is closest to **partially observed**  $Y$ .

## Low-rank matrices.

$$\hat{Y} = UV^{\top}, \quad \text{where } U \in \mathbb{R}^{n \times d}, \quad V \in \mathbb{R}^{m \times d}, \quad d \ll \min(m, n)$$



# WHY LOW-RANK?

Suppose there are  $d$  pure **types of users**,  
with rating preferences

$$V_{*1}, V_{*2}, \dots, V_{*d} \in \mathbb{R}^m.$$

Assume every user's rating can be expressed  
as a **weighted sum** of these pure ratings:

$$Y_{a*} = U_{a1} V_{*1} + U_{a2} V_{*2} + \dots + U_{ad} V_{*d}.$$

Then,  $Y = UV^\top$  where  $V_{*1}, V_{*2}, \dots, V_{*d}$  are the columns of  $V$ .

**Question.** What happens when  $d = \min(m, n)$ ?





# ALGORITHM

## Coordinate Descent.

$$\begin{aligned}\mathcal{L}(U, V) &= \sum_{(a,i) \in D} \frac{1}{2} (Y_{ai} - (UV^\top)_{ai})^2 + \frac{\lambda}{2} \|U\|^2 + \frac{\lambda}{2} \|V\|^2 \\ &= \sum_{(a,i) \in D} \frac{1}{2} (Y_{ai} - u^{(a)} \cdot v^{(i)})^2 + \frac{\lambda}{2} \sum_a \|u^{(a)}\|^2 + \frac{\lambda}{2} \sum_i \|v^{(i)}\|^2\end{aligned}$$

where  $u^{(a)}$  is the  $a$ -th row of  $U$  and  $v^{(i)}$  is the  $i$ -th row of  $V$ .

Repeat until convergence:

1. Fix  $V$  and minimize  $\mathcal{L}(U, V)$  over  $U$ .
2. Fix  $U$  and minimize  $\mathcal{L}(U, V)$  over  $V$ .



# ALGORITHM

1. Randomly initialize  $v^{(1)}, v^{(2)}, \dots, v^{(m)}$ .
2. Repeat until convergence:
  - a. For each user  $a$ , find  $u^{(a)}$  that minimizes
$$\sum_{i: (a,i) \in D} \frac{1}{2} (Y_{ai} - u^{(a)} \cdot v^{(i)})^2 + \frac{\lambda}{2} \|u^{(a)}\|^2$$
  - b. For each movie  $i$ , find  $v^{(i)}$  that minimizes
$$\sum_{a: (a,i) \in D} \frac{1}{2} (Y_{ai} - u^{(a)} \cdot v^{(i)})^2 + \frac{\lambda}{2} \|v^{(i)}\|^2$$

These are  
standard  
linear  
regression  
problems.



# DISCUSSION

## Optimization.

1. Like k-means, the algorithm converges to a local minimum.
2. Perform multiple initializations, and pick best result.

## Generalization.

1. Use **validation** to pick right **hyperparameters**  $d$  and  $\lambda$ .
  - a. Split data set into Training Data and Validation Data.
  - b. For each  $d$  and  $\lambda$ , train a predictor using Training Data.
  - c. Choose predictor that minimizes Validation Error.

